



Figure 7: The count (y-axis) of occurrences (x-axis) of each generated measure appearing in the training set.

Related Work

Automatic music generation has a long history, and recent work has focused on using artificial intelligence (?) and specifically deep neural networks. Such efforts are driven by not only the prospect of having AI aid or replace human composers and arrangers for a variety of music, but also the pursuit of probing the artistic creativity of the state-of-the-art data-driven models. Most modern work in automatic music generation leverages model architectures that have shown to be effective in computer vision or NLP, such as LSTM (?), Transformers (???), or custom architectures to learn an embedding space (?). However, the majority of the work on music generation with AI has happened in a supervised setting, which greatly limits its application to lesser known genres or specific instruments, such as drum generation. While there is yet to be a pre-trained music generation model as versatile as its counterpart in NLP such as GPT3, we believe the exploration of language-to-music transfer is necessary but uncharted.

Automatic drum generation has a small body of existing work. Similar to music generation in general, the drum generation task can happen in various settings. In the simplest setting, only one measure of drum pattern (also known as a “beat”) is generated that is supposed to repeat throughout a song (???), while we focus on a more involved setting of non-repetitive drum composition. Alternatively, some work simply considers the general rhythm (?), but not the orchestration of different drums that we emphasize on. In more practical settings, a long sequence of drum composition is generated conditioned on musical signals such as the basslines (?). While this line of work is most similar to ours which by far only deal with drum solo performance, all the work above has used drum composition data with limited size and variations as well as models (such as LSTM) that are relatively outdated in the AI community. Nevertheless, a direct comparison would be beneficial in future work. Less related is another line of work focuses on the microtiming and humanization of drum performance (???), striving to mimic human’s expressive imperfections.

A small body of work has focused on rhythm games (?). While rhythm games and drumming have certain similarities such as the focus on note placement with regard to the rhythm, their core difference lies in that the choreography of rhythm games is optimized for difficulty and playability, while the composition of drums is optimized for musicality. Due to this fundamental difference of motivation, we consider this line of work to be mostly irrelevant to ours.

Large language models (LLMs) are deep neural models, such as Transformers, pre-trained on a massive text corpus. For example, GPT3 is pre-trained on the compilation of Common Crawl, containing most texts in the world wide net, publicly available books, Wikipedia, and so on. From BERT (?) to GPT3 (?), LLMs have been dominant in most tasks and applications in NLP. While much about how LLMs work is unknown, and thus LLMs are notoriously known as black-boxes, there is a generally consensus that LLMs’ power can be attributed to the large size of both the pre-training data and the model, which give rise to LLMs’ ability to effectively adapt to low-resource domains via transfer learning by being finetuned on a small amount of data. Interestingly, a few recent work has found that some of these abilities include transfer from pre-trained language to non-language tasks, such as chess (?). While non-language textual representations such as chess moves or music charts are similar to natural language superficially, each manifests vastly different structures, and so we claim that transfer learning between them is well worth studying.

Conclusion and Future Work

Our preliminary findings show that pre-trained large language models (LLMs) finetuned with merely hundreds of symbolic music files, such as drum grooves, can learn to generate music non-trivially. We also provide evidence that such ability can be attributed both the model size and the presence of language pre-training. We hope this observation inspires research efforts in not only low-resource music generation, but also exploration of extraordinary potentials of LLMs. We also attempt to pioneer the automatic evaluation of drum grooves which we hope to facilitate future work in drum generation with AI.

We also briefly discuss our plans for future work. While drum generation is scarce in literature and challenging to reproduce, we will still strive compare some existing specialized models. While our current drumroll representation ignores velocity, such information can easily be encoded by replacing the marker ‘o’ with the velocity value. However, the effects of doing so remains to be explored. Our methodology might be ported to other instruments such as piano, which we plan to explore, whereas it would be more involved to tackle multi-instrument conditioned generation.

Porro maxime autem, corrupti excepturi quaerat veritatis odit, sunt minima iure molestias non asperiores repudian-dae inventore sint, pariat doloribus reprehenderit fugit impedit laudantium?Soluta saepe tenetur perferendis alias dolorem corporis explicabo illo ad id, temporibus eius ab, alias iusto itaque deleniti deserunt unde similique non ducimus aliquid dicta provident?Alias inventore nesciunt, harum sunt ipsam?Minus praesentium accusantium explicabo commodi veniam quis quisquam, hic rerum amet, reprehenderit inventore nobis tenetur voluptate consectetur eaque facere perferendis, unde deserunt officia est tempora sint ex placeat, vel rerum assumenda commodi aut exercitationem officiis unde illum illo accusantium?Totam expedita corrupti porro nobis ratione aliquid asperiores quasi aliquam esse eligendi, ipsa architecto amet fuga quam asperiores esse fugiat quia non vitae quidem, quisquam architecto debitis dolores quae dig-

nissimos magnam at possimus pariatur distinctio?Unde vitae
iure a architecto nemo aliquid harum totam, ut id harum